

# An implementation of the SMite virtual machine for POSIX version 0.1

Reuben Thomas

12th January 2019

## 1 Introduction

The SMite virtual machine [2] provides a portable virtual machine environment for study and experiment. To this end, SMite is itself written in ISO C, since almost all systems have an ISO C compiler available for them.

As well as the virtual machine, C SMite provides a debugger, which is described in [1].

The SMite virtual machine is described in [2]. This paper only describes the features specific to this implementation.

## 2 Using C SMite

This section describes how to compile C SMite, and the exact manner in which the interface calls and SMite's memory and registers should be accessed.

### 2.1 Configuration

SMite is written in ISO C99 using POSIX-1.2001 APIs.

The SMite virtual machine is inherently 32-bit, but will run happily on systems with larger (or smaller) addresses.

### 2.2 Compilation

SMite's build system is written with GNU autotools, and the user needs only standard POSIX utilities to run it. Installation instructions are provided in the top-level file `README.md`.

### 2.3 Registers and memory

SMite's registers are declared in `smite.h`. Their names correspond to those given in [2, section 2.1], although some have been changed to meet the rules for C identifiers. C SMite does not allocate any memory for SMite, nor does it

initialise any of the registers. C SMite provides the interface call **smite\_init()** to do this (see section 2.5).

The variables EP, I, A, MEMORY and STACK\_DEPTH correspond exactly with the SMite registers they represent, and may be read and assigned to accordingly, bearing in mind the restrictions on their use given in [2].

C SMite provides the ability to map native memory blocks into SMite's address space; see below.

An additional pseudo-register is provided by C SMite: **S0**, the address of the base of the stack.

## 2.4 Extra instructions

C SMite provides some extra action instructions, which give access to various library functions. They take the following form:

Stack effect	Description
$i*x \ u - j*x$	Call the $uth$ function of the library, passing arguments $i*x$ , with results $j*x$ .

Error  $-9$  is raised for an invalid function.

### 2.4.1 SMite

SMite provides access to its own API via the **LIB\_SMITE** instruction, opcode 0x3f. This allows the current state to be controlled, or fresh states to be created and controlled. The API is mapped to the following functions:

Code	Name	Stack effect	Description
0x0	CURRENT_STATE	$- \ n\text{-}addr$	a pointer to the current state
0x1	LOAD_WORD	$n\text{-}addr \ a\text{-}addr - x \ n$	
0x2	STORE_WORD	$n\text{-}addr \ a\text{-}addr \ x - n$	
0x3	LOAD_BYTE	$n\text{-}addr \ addr \ x \ n$	
0x4	STORE_BYTE	$n\text{-}addr \ addr \ x - n$	
0x5	REALLOC_MEMORY	$n\text{-}addr_1 \ n\text{-}addr_2 \ u \ f - addr$	
0x6	REALLOC_STACK	$n\text{-}addr_1 \ n\text{-}addr_2 \ u \ f - addr$	
0x7	NATIVE_ADDRESS_OF_RANGE	$n\text{-}addr_1 \ addr \ f - n\text{-}addr_2$	
0x8	RUN	$n\text{-}addr - n$	
0x9	SINGLE_STEP	$n\text{-}addr - n$	
0xa	LOAD_OBJECT	$n\text{-}addr \ fid \ addr - n$	
0xb	INIT	$a\text{-}addr \ u_1 \ u_2 \ u_3 - n\text{-}addr$	
0xc	DESTROY	$n\text{-}addr$	
0xd	REGISTER_ARGS	$n\text{-}addr \ n\text{-}addr \ n - n$	

### 2.4.2 Standard library

Standard C runtime and library functionality is accessed via the **LIB\_C** instruction, opcode 0x3e.

### 2.4.2.1 Command-line arguments

Some functions are provided to access command-line arguments passed to C SMite (excluding any that it interprets itself).

Code	Name	Stack effect	Description
0x0	ARGC	- u	the number of arguments
0x1	ARG_LEN	$u_1 - u_2$	the length of the $u_1$ th argument
0x2	ARG_COPY	$u_1$ c-addr $u_2 - u_3$	copy argument $u_1$ to the buffer of length $u_2$ at c-addr, leaving t

### 2.4.2.2 Standard I/O streams

These extra instructions provide access to POSIX standard input, output and error. Each call returns a corresponding file identifier.

Opcode	POSIX file descriptor
0x3	STDIN_FILENO
0x4	STDOUT_FILENO
0x5	STDERR_FILENO

### 2.4.2.3 File system

Opcode	Forth word
0x6	OPEN-FILE
0x7	CLOSE-FILE
0x8	READ-FILE
0x9	WRITE-FILE
0xa	FILE-POSITION
0xb	REPOSITION-FILE
0xc	FLUSH-FILE
0xd	RENAME-FILE
0xe	DELETE-FILE
0xf	FILE-SIZE
0x10	RESIZE-FILE
0x11	FILE-STATUS

The implementation-dependent word returned by FILE-STATUS contains the POSIX protection bits, given by the `st_mode` member of the struct `stat` returned for the given file descriptor.

File access methods are bit-masks, composed as follows:

Bit value	Meaning
1	read
2	write
4	binary mode

To create a file, set both read and write bits to zero when calling OPEN-FILE.

## 2.5 Using the interface calls

The operation of the specified interface calls is given in [2]. Here, the C prototypes corresponding to the idealised prototypes used in [2] are given. The names are prefixed with **smite\_**. The first argument to most routines is a `@PACKAGE_state *`, as returned by `smite_init`.

```
uint8_t *native_address_of_range(smite_state *state,
smite_UWORD address, smite_UWORD length)
```

Returns NULL when the address range is not entirely valid.

```
smite_state *state, smite_WORD run(smite_state *state)
```

The reason code returned by **smite\_run()** is a SMite word.

```
smite_state *state, smite_WORD smite_single_step(smite_state
*state)
```

The reason code returned by **smite\_single\_step()** is a SMite word.

```
ptrdiff_t smite_load_object(smite_state *state, smite_UWORD
address, int fd)
```

If a file system error occurs, the return code is `-3`. If the endism of the object file does not match `ENDISM`, the return code is `-4`. If the word size of the object file is not `WORD_SIZE`, the return code is `-5`. As an extension to the specification, if an object file starts with the bytes `35, 33 (!)`, then it is assumed to be the start of a UNIX-style “hash bang” line, and the file contents up to and including the first newline character (10) is ignored.

In addition to the required interface calls C SMite provides an initialisation routine **smite\_init()** which, given a word array and its size, initialises SMite:

```
smite_state *smite_init(size_t memory_size, size_t stack_size)
memory_size is the size of b_array in words (not bytes); similarly,
stack_size gives the size of the stack in words; these are allocated by
smite_init. The return value is NULL if memory cannot be allocated,
or if any requested size is too large, and a pointer to a new state
otherwise. All the registers are initialised as per [2].
```

```
int smite_realloc_memory(smite_state *state, smite_UWORD size)
Resize the memory to the given size. Any new memory is zeroed.
Returns 0 on success or -1 if the requested size of memory cannot be
allocated.
```

```
int smite_realloc_stack(smite_state *state, smite_UWORD size)
Resize the stack to the given size. Any new memory is zeroed.
Returns 0 on success or -1 if the requested size of memory cannot be
allocated.
```

The following routines give easier access to SMite’s address space at the byte and word level. On success, they return 0, and on failure, the relevant error code.

```
int smite_load_word(smite_state *state, smite_UWORD address,
smite_WORD *value)
```

Load the word at the given address into the given `smite_WORD *`.

```
int smite_store_word(smite_state *state, smite_UWORD address,
smite_WORD value)
```

Store the given WORD value at the given address.

```
int smite_load_byte(smite_state *state, smite_UWORD address,
smite_BYTE *value)
```

Load the byte at the given address into the given smite\_BYTE \*.

```
int smite_store_byte(smite_state *state, smite_UWORD address,
smite_BYTE value)
```

Store the given smite\_BYTE value at the given address.

The following routines give access to SMite's stacks. On success, they return 0, and on failure, the relevant error code.

```
int smite_load_stack(smite_state *state, smite_UWORD pos,
smite_WORD *value)
```

Load the word at the given position of the data stack into the given smite\_WORD \*.

```
int smite_store_stack(smite_state *state, smite_UWORD pos,
smite_WORD value)
```

Store the given WORD value at the given position in the data stack.

```
int smite_pop_stack(smite_state *state, smite_WORD *value)
```

Pop the top word off the data stack, decrementing its depth, into the given smite\_WORD \*.

```
int smite_push_stack(smite_state *state, smite_WORD value)
```

Push the given smite\_WORD value on to the data stack of the given size, incrementing its depth.

The following routine allows the calling program to register command-line arguments that can be retrieved by the ARGV and ARG\_COPY extra instructions.

```
int smite_register_args(smite_state *state, int argc, char
*argv[])
```

Maps the given arguments register, which has the same format as that supplied to **main()**, into SMite's memory. Returns 0 on success and -1 if memory could not be allocated, or -2 if an argument could not be mapped to SMite's address space.

Programs which use C SMite's interface must **#include** the header file `smite/smite.h` and be linked with the SMite library. `smite/opcodes.h`, which contains an enumeration type of SMite's instruction set, and `smite/smite_aux.h`, which contains other useful declarations, may also be useful; they are not documented here.

## 2.6 Other extras provided by C SMite

C SMite provides various useful extras in `smite_aux.h`. These are used internally, and are thought to be useful, but may change at any time, so their stability should not be relied on.

## References

- [1] Reuben Thomas. A simple shell for the SMite Forth virtual machine, 2018. <https://rrt.sc3d.org/>.
- [2] Reuben Thomas. The SMite Forth virtual machine, 2018. <https://rrt.sc3d.org/>.