

An implementation of the SMite virtual machine for POSIX version 0.1

Reuben Thomas

12th January 2019

1 Introduction

The SMite virtual machine [4] provides a portable virtual machine environment for study and experiment. To this end, SMite is itself written in ISO C, since almost all systems have an ISO C compiler available for them.

As well as the virtual machine, C SMite provides a debugger, which is described in [3].

The SMite virtual machine is described in [4]. This paper only describes the features specific to this implementation.

2 Using C SMite

This section describes how to compile C SMite, and the exact manner in which the interface calls and SMite's memory and registers should be accessed.

2.1 Configuration

SMite is written in ISO C99 using POSIX-1.2001 APIs.

The SMite virtual machine is inherently 32-bit, but will run happily on systems with larger (or smaller) addresses.

2.2 Compilation

SMite's build system is written with GNU autotools, and the user needs only standard POSIX utilities to run it. Installation instructions are provided in the top-level file `README.md`.

2.3 Registers and memory

SMite's registers are declared in `smite.h`. Their names correspond to those given in [4, section 2.1], although some have been changed to meet the rules for C identifiers. C SMite does not allocate any memory for SMite, nor does it

initialise any of the registers. C SMite provides the interface call **smite_init()** to do this (see section 2.5).

The variables *EP*, *I*, *A*, *MEMORY*, *SDEPTH* and *RDEPTH* correspond exactly with the SMite registers they represent, and may be read and assigned to accordingly, bearing in mind the restrictions on their use given in [4].

C SMite provides the ability to map native memory blocks into SMite's address space; see below.

Two additional pseudo-registers are provided by C SMite: they are called *S0* and *R0*, and are the address of the base of the data and return stacks respectively.

2.4 Extra instructions

C SMite provides some libraries of EXTRA instruction functions, which are called in general as:

Stack effect	Description
$i*x \ u1 \ u2 \ - \ j*x$	Call the $u1$ th function of library $u2$, passing arguments $i*x$, with results $j*x$.

Error -258 is raised if an invalid library is accessed, and -259 for an invalid function in a known library.

2.4.1 SMite

SMite provides access to its own API as library -1 . This allows the current state to be controlled, or fresh states to be created and controlled.

Code	Name	Stack effect	Description
0x0	CURRENT_STATE	$- \ n\text{-addr}$	a pointer to the current state
0x1	LOAD_WORD	$n\text{-addr} \ a\text{-addr} \ - \ x \ n$	
0x2	STORE_WORD	$n\text{-addr} \ a\text{-addr} \ x \ - \ n$	
0x3	LOAD_BYTE	$n\text{-addr} \ addr \ x \ n$	
0x4	STORE_BYTE	$n\text{-addr} \ addr \ x \ - \ n$	
0x5	MEM_REALLOC	$n\text{-addr}_1 \ n\text{-addr}_2 \ u \ f \ - \ addr$	
0x6	NATIVE_ADDRESS	$n\text{-addr}_1 \ addr \ f \ - \ n\text{-addr}_2$	
0x7	RUN	$n\text{-addr} \ - \ n$	
0x8	SINGLE_STEP	$n\text{-addr} \ - \ n$	
0x9	LOAD_OBJECT	$n\text{-addr} \ fid \ addr \ - \ n$	
0xa	INIT	$a\text{-addr} \ u_1 \ u_2 \ u_3 \ - \ n\text{-addr}$	
0xb	DESTROY	$n\text{-addr}$	
0xc	REGISTER_ARGS	$n\text{-addr} \ n\text{-addr} \ n \ - \ n$	

2.4.2 Standard library

Standard C runtime and library functionality is accessed via library 0.

2.4.2.1 Command-line arguments

Two calls are provided to access command-line arguments passed to C SMite (excluding any that it interprets itself). They are copied from Gforth ([2]).

Code	Name	Stack effect	Description
0x0	ARGC	- u	the number of arguments
0x1	ARG_LEN	u1 - u2	the length of the u1th argument
0x2	ARG_COPY	u ₁ c-addr u ₂ - u ₃	copy argument u ₁ to the buffer of length u ₂ at c-addr, leaving t

2.4.2.2 Standard I/O streams

These extra instructions provide access to POSIX standard input, output and error. Each call returns a corresponding file identifier.

Opcode	POSIX file descriptor
0x3	STDIN_FILENO
0x4	STDOUT_FILENO
0x5	STDERR_FILENO

2.4.2.3 File system

The file system extra instructions correspond directly to ANS Forth words, as defined in [1].

Opcode	Forth word
0x6	OPEN-FILE
0x7	CLOSE-FILE
0x8	READ-FILE
0x9	WRITE-FILE
0xa	FILE-POSITION
0xb	REPOSITION-FILE
0xc	FLUSH-FILE
0xd	RENAME-FILE
0xe	DELETE-FILE
0xf	FILE-SIZE
0x10	RESIZE-FILE
0x11	FILE-STATUS

The implementation-dependent word returned by FILE-STATUS contains the POSIX protection bits, given by the `st_mode` member of the struct `stat` returned for the given file descriptor.

File access methods are bit-masks, composed as follows:

Bit value	Meaning
1	read
2	write
4	binary mode

To create a file, set both read and write bits to zero when calling OPEN-FILE.

2.5 Using the interface calls

The operation of the specified interface calls is given in [4]. Here, the C prototypes corresponding to the idealised prototypes used in [4] are given. The names are prefixed with **smite_**. The first argument to most routines is a `@PACKAGE_state *`, as returned by `smite_init`.

```
uint8_t *native_address_of_range(smite_state *state,  
smite_UWORD address, smite_UWORD length)
```

Returns NULL when the address range is not entirely valid.

```
smite_state *state, smite_WORD run(smite_state *state)
```

The reason code returned by **smite_run()** is a SMite word.

```
smite_state *state, smite_WORD smite_single_step(smite_state  
*state)
```

The reason code returned by **smite_single_step()** is a SMite word.

```
ptrdiff_t smite_load_object(smite_state *state, smite_UWORD  
address, int fd)
```

If a file system error occurs, the return code is `-3`. If the endism of the object file does not match `ENDISM`, the return code is `-4`. If the word size of the object file is not `WORD_SIZE`, the return code is `-5`. As an extension to the specification, if an object file starts with the bytes 35, 33 (!), then it is assumed to be the start of a UNIX-style “hash bang” line, and the file contents up to and including the first newline character (10) is ignored.

In addition to the required interface calls C SMite provides an initialisation routine **smite_init()** which, given a word array and its size, initialises SMite:

```
smite_state *smite_init(size_t memory_size, size_t  
data_stack_size, size_t return_stack_size)
```

`memory_size` is the size of `b_array` in *words* (not bytes); similarly, `data_stack_size` and `return_stack_size` give the size of the stacks in words; these are allocated by `smite_init`. The return value is NULL if memory cannot be allocated, or if any requested size is too large, and a pointer to a new state otherwise. All the registers are initialised as per [4].

```
int smite_mem_realloc(smite_state *state, smite_UWORD size)
```

Resize the memory to the given size. Any new memory is zeroed. Returns 0 on success or `-1` if the requested size of memory cannot be allocated.

The following routines give easier access to SMite’s address space at the byte and word level. On success, they return 0, and on failure, the relevant error code.

```
int smite_load_word(smite_state *state, smite_UWORD address,  
smite_WORD *value)
```

Load the word at the given address into the given `smite_WORD *`.

```
int smite_store_word(smite_state *state, smite_UWORD address,
smite_WORD value)
```

Store the given WORD value at the given address.

```
int smite_load_byte(smite_state *state, smite_UWORD address,
smite_BYTE *value)
```

Load the byte at the given address into the given smite_BYTE *.

```
int smite_store_byte(smite_state *state, smite_UWORD address,
smite_BYTE value)
```

Store the given smite_BYTE value at the given address.

The following routines give access to SMite's stacks. On success, they return 0, and on failure, the relevant error code.

```
int smite_load_stack(smite_state *state, smite_UWORD pos,
smite_WORD *value)
```

Load the word at the given position of the data stack into the given smite_WORD *.

```
int smite_store_stack(smite_state *state, smite_UWORD pos,
smite_WORD value)
```

Store the given WORD value at the given position in the data stack.

```
int smite_pop_stack(smite_state *state, smite_WORD *value)
```

Pop the top word off the data stack, decrementing its depth, into the given smite_WORD *.

```
int smite_push_stack(smite_state *state, smite_WORD value)
```

Push the given smite_WORD value on to the data stack of the given size, incrementing its depth.

```
int smite_load_return_stack(smite_state *state, smite_UWORD
pos, smite_WORD *value)
```

Load the word at the given position of the return stack into the given smite_WORD *.

```
int smite_store_return_stack(smite_state *state, smite_UWORD
pos, smite_WORD value)
```

Store the given WORD value at the given position in the return stack.

```
int smite_pop_return_stack(smite_state *state, smite_WORD
*value)
```

Pop the top word off the return stack, decrementing its depth, into the given smite_WORD *.

```
int smite_push_return_stack(smite_state *state, smite_WORD
value)
```

Push the given smite_WORD value on to the return stack of the given size, incrementing its depth.

The following routine allows the calling program to register command-line arguments that can be retrieved by the ARGV and ARG_COPY extra instructions.

```
int smite_register_args(smite_state *state, int argc, char
*argv[])
```

Maps the given arguments register, which has the same format as that supplied to **main()**, into SMite’s memory. Returns 0 on success and -1 if memory could not be allocated, or -2 if an argument could not be mapped to SMite’s address space.

The following routine allows object files to be saved:

```
int smite_save_object(smite_state, *statesmite_UWORD address,
smite_UWORD length, int fd)
```

On success, 0 is returned. If the address or length are invalid, the return code is -1 . If a file system error occurs, the return code is -2 .

Programs which use C SMite’s interface must `#include` the header file `smite.h` and be linked with the SMite library. `smite_opcodes.h`, which contains an enumeration type of SMite’s instruction set, and `smite_debug.h`, which contains useful debugging functions such as disassembly, may also be useful; they are not documented here.

2.6 Other extras provided by C SMite

C SMite provides various useful extras in `smite_aux.h`. These are used internally, and are thought to be useful, but may change at any time, so their stability should not be relied on.

References

- [1] American National Standards Institute. *ANS X3.215-1994: Programming Languages—Forth*, 1994.
- [2] M. Anton Ertl, Bernd Paysan, Jens Wilke, Neal Crook, David Kühling, et al. Gforth. <https://www.complang.tuwien.ac.at/forth/gforth/>, 1993.
- [3] Reuben Thomas. A simple shell for the SMite Forth virtual machine, 2018. <https://rrt.sc3d.org/>.
- [4] Reuben Thomas. The SMite Forth virtual machine, 2018. <https://rrt.sc3d.org/>.