The Mit Virtual Machine

Reuben Thomas

28th March 2020

Typographical conventions

Instructions and registers are shown in `typewriter` font; interface calls are shown in **bold** type.

Addresses are given in bytes and refer to the VM address space except where stated. Addresses are written in hexadecimal; hex numbers are prefixed with "0x".

1   Introduction

Mit is a simple virtual machine for study and experiment. It is a stack machine, based on the more complex register machine Mite [4]. This paper gives a full description of Mit.

Mit is conceptually (and usually in fact) a library, embedded in other programs.

2   Architecture

The address unit is the byte, which is eight bits. Most of the quantities on which Mit operates are fixed-size words, which are stored in memory in either big- or little-endian order. The choice of byte and word size enable efficient implementation on the vast majority of machine architectures.

2.1   Parameters

Some parameters may take a range of values, but are fixed for a particular virtual machine instance. They are listed in table 1.

Parameter	Meaning
`word_bytes`	The number of bytes in a word. Must be in the range 2 to 32 inclusive, and a power of 2.
`instruction_bits`	The number of bits in an instruction. Must be in the range 5 to 8 inclusive.

Table 1: Parameters

2.2 Registers

The registers are word quantities; they are listed, with their functions, in table 2. The registers are initialised to 0.

Register	Function
pc	The program counter. Points to the next word from which i may be loaded.
ir	The instruction register. Contains instructions to be executed.
stack_depth	The number of words on the stack.

Table 2: Registers

2.3 Memory

Mit's memory consists of discontiguous words in a single flat address space. The address of a word is that of the byte in it with the lowest address.

2.4 Stack

The stack is a LIFO stack of words used for passing values to instructions and routines and for holding subroutine return addresses. To **push** a word on to the stack means to add a new word to the top of the stack, increasing the stack depth by 1; to **pop** a word means to reduce the stack depth by 1. Instructions that change the number of words on the stack implicitly pop their arguments and push their results.

2.5 Execution

Execution proceeds as follows:

> begin
>> let **opcode** be the least significant instruction_bits bits of ir
>> shift ir logically instruction_bits bits to the right
>> execute the instruction given by **opcode**
> repeat

If an error occurs during execution (see section 2.6), the state of the virtual machine is reset to its state at the start of the loop before the error is raised. This allows instructions to be restarted after handling the error, where desired.

2.6 Errors and termination

When Mit encounters certain abnormal situations, such as an attempt to access an invalid address, or divide by zero, an **error** is **raised**, and execution terminates. The effect of the current instruction is undone (see section 2.5). An **error code** is returned to the caller.

Execution can be terminated explicitly by performing a `halt` instruction (see section 3.3.1).

Error codes are unsigned numbers. 0 to 127 are reserved for the specification; other error codes may be used by implementations. The meanings of those that may be raised by Mit are shown in table 3.

Code	Meaning
0	**single_step()** has terminated without error.
1	Invalid opcode (see section 3.11).
2	Stack overflow.
3	Invalid stack read.
4	Invalid stack write.
5	Invalid memory read.
6	Invalid memory write.
7	Address alignment error: raised when an instruction is given a valid address, but insufficiently aligned.
8	Invalid size (greater than $\log_2$ `word_bytes`).
9	Division by zero attempted (see section 3.8).
127	A `halt` instruction was executed.

Table 3: Errors raised by Mit

3 Instruction set

The instruction set is listed below, with the instructions grouped according to function. The instructions are given in the following format:

`NAME`   (*before* - *after*)
Description.

The first line consists of the name of the instruction. On the right is the stack effect, which shows the effect of the instruction on the stack. Underneath is the description.

Stack effects are written

(*before* - *after*)

where *before* and *after* are stack pictures showing the items on top of a stack before and after the instruction is executed. An instruction only affects the items shown in its stack effects. The brackets and dashes serve merely to delimit the stack effect and to separate *before* from *after*. **Stack pictures** are a representation of the top-most items on the stack, and are written

i_1 i_2...i_{n-1} i_n

where the i_k are stack items, each of which occupies a whole number of words,[1] with i_n being on top of the stack. The symbols denoting different types of stack item are shown in table 4.

Types are only used to indicate how instructions treat their arguments and results; Mit does not distinguish between stack items of different types. In stack

---

[1] In this specification, each stack item occupies *precisely* one word.

Symbol	Data type
flag	a Boolean flag, 0 for false or non-zero for true
size	an integer in the range 0 to $\log_2$ `word_bytes` inclusive
n	signed number
u	unsigned number
n|u	number (signed or unsigned)
x	unspecified word
addr	address
a-addr	word-aligned address

Table 4: Types used in stack effects

pictures the most general argument types with which each instruction can be supplied are given; subtypes may be substituted. Using the phrase "$i \Rightarrow j$" to denote "i is a subtype of j", table 5 shows the subtype relationships. The subtype relation is transitive.

$$u \Rightarrow x$$
$$n \Rightarrow x$$
$$flag \Rightarrow u$$
$$size \Rightarrow u$$
$$a\text{-}addr \Rightarrow addr \Rightarrow u$$

Table 5: The subtype relation

Numbers are represented in two's complement form. *addr* consists of all valid virtual machine addresses.

Each type may be suffixed by a number in stack pictures; if the same combination of type and suffix appears more than once in a stack effect, it refers each time to the identical stack item.

Ellipsis is used for indeterminate numbers of specified types of item.

3.1 Instruction fetch

If an invalid or unaligned address is accessed when loading `ir`, the appropriate error is raised (see section 2.6).

`next` (-)

Load the word pointed to by `pc` into `ir` then add `word_bytes` to `pc`.

3.2 Control

These instructions implement unconditional and conditional branches, and subroutine call and return (subroutine return is `jump`):

`jump` (*a-addr* -)

If `ir` is not 0, raise error 1. Set `pc` to *a-addr*. Perform the action of `next`.

```
jumpz                                          ( flag a-addr - )
```
If *flag* is false then set `pc` to *a-addr* and perform the action of `next`.

```
call                                      ( a-addr₁ - a-addr₂ )
```
If `ir` is not 0, raise error 1. Exchange `pc` with the top stack value. Perform the action of `next`.

3.3   Extra instructions

Since `ir` must be 0 when `next` is performed, the rest of an instruction word following `next`, `jump` and `call` must normally be all zero bits.

Non-zero values following `next` and `call` are reserved for the Mit specification; non-zero values following `jump` may be used by implementations to implement extra functionality.

When an extra instruction is performed, the original instruction is considered to have completed executing.

3.3.1   Termination

This instruction terminates execution (see section 2.6):

```
halt                                                      ( - )
```
Raise error 127.

3.4   Stack manipulation

These instructions manage the stack:

```
pop                                                       ( x - )
```
Remove *x* from the stack.

```
dup                               ( xᵤ...x₀ u - xᵤ...x₀ xᵤ )
```
Remove *u*. Copy x_u to the top of the stack.

```
swap                       ( x_{u+1}...x₀ u - x₀ xᵤ...x₁ x_{u+1} )
```
Exchange the top stack word with the *u+1*th.

3.5   Literals

```
lit                                                       ( - n )
```
The word pointed to by `pc` is pushed on to the stack, and `pc` is incremented to point to the following word.

```
lit_pc_rel                                                ( - n )
```
Like `lit`, except that the initial value of `pc` is added to the value pushed on to the stack.

```
lit_0                                                     ( - 0 )
```
Push 0 on to the stack.

```
lit_1                                                     ( - 1 )
```
Push 1 on to the stack.

```
lit_2                                                     ( - 2 )
```
Push 2 on to the stack.

```
lit_3                                                            ( - 3 )
```
Push 3 on to the stack.

3.6 Logic and shifts

Logic functions:

```
not                                                        ( x₁ - x₂ )
```
Invert all bits of x_1, giving its logical inverse x_2.

```
and                                                    ( x₁ x₂ - x₃ )
```
x_3 is the bit-by-bit logical "and" of x_1 with x_2.

```
or                                                     ( x₁ x₂ - x₃ )
```
x_3 is the bit-by-bit inclusive-or of x_1 with x_2.

```
xor                                                    ( x₁ x₂ - x₃ )
```
x_3 is the bit-by-bit exclusive-or of x_1 with x_2.

Shifts:

```
lshift                                                  ( x₁ u - x₂ )
```
Perform a logical left shift of u bit-places on x_1, giving x_2. Put zero into the least significant bits vacated by the shift. If u is greater than or equal to the number of bits in a word, x_2 is zero.

```
rshift                                                  ( x₁ u - x₂ )
```
Perform a logical right shift of u bit-places on x_1, giving x_2. Put zero into the most significant bits vacated by the shift. If u is greater than or equal to the number of bits in a word, x_2 is zero.

```
arshift                                                 ( x₁ u - x₂ )
```
Perform an arithmetic right shift of u bit-places on x_1, giving x_2. Copy the original most-significant bits into the most significant bits vacated by the shift. If u is greater than or equal to the number of bits in a word, all the bits of x_2 are the same as the original most-significant bit.

```
sign_extend                                           ( u size - n )
```
Sign extend the 2^{size}-byte quantity u to n.

3.7 Comparison

These words compare two numbers on the stack, returning a flag (for equality, use xor; see section 3.6):

```
lt                                                   ( n₁ n₂ - flag )
```
flag is 1 if and only if n_1 is less than n_2.

```
ult                                                  ( u₁ u₂ - flag )
```
flag is 1 if and only if u_1 is less than u_2.

3.8 Arithmetic

These instructions consist of monadic and dyadic operators. All calculations are made without bounds or overflow checking, except as detailed for certain instructions.

Negation and addition:

`negate`   (n_1 – n_2)

Negate n_1, giving its arithmetic inverse n_2.

`add`   (n_1|u_1 n_2|u_2 – n_3|u_3)

Add n_2|u_2 to n_1|u_1, giving the sum n_3|u_3.

Multiplication and division (note that all division instructions raise error 9 if division by zero is attempted):

`mul`   (n_1|u_1 n_2|u_2 – n_3|u_3)

Multiply n_1|u_1 by n_2|u_2 giving the product n_3|u_3.

`divmod`   (n_1 n_2 – n_3 n_4)

Divide n_1 by n_2 using symmetric division, giving the single-word quotient n_3 and the single-word remainder n_4. The quotient is rounded towards zero.

`udivmod`   (u_1 u_2 – u_3 u_4)

Divide u_1 by u_2, giving the single-word quotient u_3 and the single-word remainder u_4.

3.9 Memory

These instructions fetch and store quantities to and from memory. If an invalid or unaligned address is accessed, the appropriate error is raised (see section 2.6).

`load`   (*addr size* – *x*)

Load the 2^{size}-byte quantity *x* stored at *addr*, which must be a multiple of 2^{size}. Any unused high-order bits are set to zero.

`store`   (*x addr size* –)

Store the 2^{size} least-significant bytes of *x* at *addr*, which must be a multiple of 2^{size}.

3.10 Instruction encoding

Instructions are encoded as `instruction_bits`-bit opcodes; opcodes are packed into words, which are executed starting at the least-significant bits.

3.11 Instruction opcodes

Table 6 lists the instruction opcodes in numerical order. Table 7 lists the extra instruction opcodes (following `call`; see section 3.3). Other instruction opcodes are undefined.

Opcode	Instruction	Opcode	Instruction
0x0	next	0x10	(undefined)
0x1	jump	0x11	lt
0x2	jumpz	0x12	ult
0x3	call	0x13	negate
0x4	pop	0x14	add
0x5	dup	0x15	mul
0x6	swap	0x16	divmod
0x7	(undefined)	0x17	udivmod
0x8	load	0x18	not
0x9	store	0x19	and
0xa	lit	0x1a	or
0xb	lit_pc_rel	0x1b	xor
0xc	lit_0	0x1c	lshift
0xd	lit_1	0x1d	rshift
0xe	lit_2	0x1e	arshift
0xf	lit_3	0x1f	sign_extend

Table 6: Instruction opcodes

Opcode	Instruction
0x1	halt

Table 7: Extra instruction opcodes

4   External interface

- Implementations should provide an **API** to create and run virtual machine instances, and provide access to its registers, stack and memory.

- Implementations can add **extra instructions** to provide extra computational primitives, and to offer access to system facilities, previously written code, native libraries and so on.

Acknowledgements

References

[1] Paulo Bonzini. Using and porting GNU *lightning*, 2000. `ftp://alpha.gnu.org/gnu/`.

[2] Martin Richards. Cintcode distribution, 2000. `https://www.cl.cam.ac.uk/~mr/BCPL.html`.

[3] Reuben Thomas. Beetle and pForth: a Forth virtual machine and compiler. BA dissertation, University of Cambridge, 1995. `https://rrt.sc3d.org/`.

[4] Reuben Thomas. *Mite: a basis for ubiquitous virtual machines*. PhD thesis, University of Cambridge Computer Laboratory, November 2000. `https://rrt.sc3d.org/`.