

Mit virtual machine specification

Reuben Thomas

20th July 2020

1 Introduction

Mit is a simple virtual machine, designed to be both easy to implement efficiently on most widely-used hardware, and to compile for. It aims to be formally tractable. This specification is intended for those who wish to implement or program Mit.

2 Parameters

The virtual machine has the following parameters:

| | |
|------------|---|
| Endianness | Memory can be either big- or little-endian. |
| word_bytes | The number of bytes in a word, 4 or 8. |

3 Memory

The flat linear address space contains `word_bytes`-byte words of 8-bit bytes. Addresses range from 0 to $2^{8 \times \text{word_bytes}} - 1$ inclusive and identify a byte; the address of a quantity is that of the byte in it with the lowest address. Whether a given word may be read or written can change during execution.

4 Registers

The registers are word quantities:

| Register | Function |
|----------|--|
| pc | The program counter. Points to the next word from which <code>ir</code> may be loaded. |
| ir | The instruction register. Contains instructions to be executed. |

5 Computation stack

Computation is performed with a last in, first out stack of words. The computation stack is usually referred to simply as “the stack”. To **push** a word on to the

stack means to add a new word to the top; to **pop** a word means to discard the top item. Instructions implicitly pop their arguments and push their results.

Stack effects are written

$$before \rightarrow after$$

where *before* and *after* are stack pictures showing the items on top of the stack before and after the instruction is executed. An instruction only affects the items shown in its stack effect. A stack item $[x]$ in square brackets is optional.

Stack pictures represent the topmost stack items, and are written

$$i_n \ i_{n-1} \dots i_2 \ i_1$$

where the i_k are stack items, with i_1 being on top of the stack.

6 Call stack and catch stack

Subroutines are implemented with the call stack, a last in, first out stack of computation stacks. The top-most is the current computation stack, which is used by instructions.

When a subroutine call is performed, a new computation stack is pushed on to the call stack. This is shown in a stack effect as:

$$caller \ ; \ callee$$

where *caller* and *callee* are stack pictures for the caller and callee respectively.

Error handling is implemented with the catch stack, a last in, first out stack of call stacks. The top-most is the current call stack. This is shown in a stack effect as:

$$handler \ | \ handlee$$

where *handler* and *handlee* are stack pictures for the top-most computation stacks in adjacent call stacks.

7 Execution

The registers (see section 4) are initialised to desired values. Execution proceeds as follows:

```
repeat
  let opcode be the least significant byte of ir
  shift ir arithmetically one byte to the right
  execute the instruction given by opcode,
  or throw error  $-1$  if the opcode is invalid
```

Instruction fetch means setting **ir** to the word pointed to by **pc** and making **pc** point to the next word. This is done whenever **ir** is 0 or -1 . These are encoded respectively as extra instruction 0 (see section 8.9) and trap -1 (see section 8.10).

7.1 Errors and termination

In exceptional situations, such as an invalid memory access, or division by zero, an **error** may be **thrown**; see section 8.8. An **error code** is returned to the handler.

Execution can be terminated explicitly by a **throw** instruction (see section 8.8), which throws an error.

Error codes are signed numbers. 0 to −127 are reserved for the specification. The specified error codes are:

| Code | Meaning |
|------|---|
| 0 | Execution has terminated without error. |
| −1 | Invalid opcode (see section 8.11). |
| −2 | Stack overflow. |
| −3 | Invalid stack read. |
| −4 | Invalid stack write. |
| −5 | Invalid memory read. |
| −6 | Invalid memory write. |
| −7 | Address is valid but insufficiently aligned. |
| −8 | Division by zero attempted (see section 8.4). |
| −9 | Division overflow (see section 8.4). |

Errors −2 to −9 inclusive are optional: an implementation may choose to raise them, or not.

8 Instructions

The instructions are listed below, grouped according to function, in the following format:

NAME *before* → *after*
Description.

The first line consists of the name of the instruction on the left, and the stack effect on the right. Underneath is the description.

The symbols denoting different types of stack item are shown in table 1.

| Symbol | Data type |
|---------------|---|
| <i>flag</i> | a Boolean flag, zero for false or non-zero for true |
| <i>s</i> | signed number |
| <i>u</i> | unsigned number |
| <i>n</i> | number (signed or unsigned) |
| <i>x</i> | unspecified word |
| <i>addr</i> | address |
| <i>a-addr</i> | word-aligned address |

Table 1: Types used in stack effects

Numbers are represented in two's complement form. *addr* consists of all valid virtual machine addresses.

Each type may be suffixed by a number in stack pictures; if the same combination of type and suffix appears more than once in a stack effect, it refers each time to the identical stack item.

Ellipsis is used for indeterminate numbers of specified types of item.

8.1 Stack manipulation

These instructions manage the stack:

| | |
|---|---|
| pop | $x \rightarrow$ |
| Remove x from the stack. | |
| dup | $x_u \dots x_0 \ u \rightarrow x_u \dots x_0 \ x_u$ |
| Remove u . Copy x_u to the top of the stack. | |
| set | $x_{u+1} \dots x_0 \ u \rightarrow x_0 \ x_u \dots x_1$ |
| Set the $u+1$ th stack word to x_0 , then pop x_0 . | |
| swap | $x_{u+1} \dots x_0 \ u \rightarrow x_0 \ x_u \dots x_1 \ x_{u+1}$ |
| Exchange the top stack word with the $u+1$ th. | |

8.2 Memory

These instructions fetch and store quantities to and from memory.

| | |
|--|---------------------------------|
| load | $a\text{-}addr \rightarrow x$ |
| Load the word x stored at $a\text{-}addr$. | |
| store | $x \ a\text{-}addr \rightarrow$ |
| Store x at $a\text{-}addr$. | |
| load1 | $addr \rightarrow x$ |
| Load the byte x stored at $addr$. Unused high-order bits are set to zero. | |
| store1 | $x \ addr \rightarrow$ |
| Store the least-significant byte of x at $addr$. | |
| load2 | $addr \rightarrow x$ |
| Load the 2-byte quantity x stored at $addr$, which must be a multiple of 2. Unused high-order bits are set to zero. | |
| store2 | $x \ addr \rightarrow$ |
| Store the 2 least-significant bytes of x at $addr$, which must be a multiple of 2. | |
| load4 | $addr \rightarrow x$ |
| Load the 4-byte quantity x stored at $addr$, which must be a multiple of 4. Any unused high-order bits are set to zero. | |
| store4 | $x \ addr \rightarrow$ |
| Store the 4 least-significant bytes of x at $addr$, which must be a multiple of 4. | |

8.3 Constants

| | |
|---|-----------------|
| push | $\rightarrow n$ |
| Push the word pointed to by pc on to the stack, and increment pc to point to the following word. | |
| pushrel | $\rightarrow n$ |
| Like push but add pc to the value pushed on to the stack. | |
| pushi_n | $\rightarrow n$ |
| Push n on to the stack. n ranges from -32 to 31 inclusive. | |
| pushreli_n | $\rightarrow n$ |
| Push $pc + \text{word_bytes} \times n$ on to the stack. n ranges from -64 to 63 inclusive. | |

The operand of pushi and pushreli is encoded in the instruction opcode; see section 8.11.

8.4 Arithmetic

All calculations are made modulo $2^{(8 \times \text{word_bytes})}$, except as detailed for certain instructions.

| | |
|--|-----------------------------------|
| neg | $s_1 \rightarrow s_2$ |
| Negate s_1 , giving s_2 . | |
| add | $n_1 \ n_2 \rightarrow n_3$ |
| Add n_2 to n_1 , giving the sum n_3 . | |
| mul | $n_1 \ n_2 \rightarrow n_3$ |
| Multiply n_1 by n_2 giving the product n_3 . | |
| divmod | $s_1 \ s_2 \rightarrow s_3 \ s_4$ |
| Divide s_1 by s_2 , giving the single-word quotient s_3 and the single-word remainder s_4 . The quotient is rounded towards zero. If s_2 is zero, throw error -8 . If s_1 is $-2^{(8 \times \text{word_bytes} - 1)}$ and s_2 is -1 , throw error -9 . | |
| udivmod | $u_1 \ u_2 \rightarrow u_3 \ u_4$ |
| Divide u_1 by u_2 , giving the single-word quotient u_3 and the single-word remainder u_4 . If u_2 is zero, throw error -8 . | |

8.5 Logic

Logic functions:

| | |
|---|-----------------------------|
| not | $x_1 \rightarrow x_2$ |
| Invert all bits of x_1 , giving its logical inverse x_2 . | |
| and | $x_1 \ x_2 \rightarrow x_3$ |
| x_3 is the bit-by-bit logical “and” of x_1 with x_2 . | |
| or | $x_1 \ x_2 \rightarrow x_3$ |
| x_3 is the bit-by-bit inclusive-or of x_1 with x_2 . | |
| xor | $x_1 \ x_2 \rightarrow x_3$ |
| x_3 is the bit-by-bit exclusive-or of x_1 with x_2 . | |

8.6 Shifts

lshift

$$x_1 \ u \rightarrow x_2$$

Perform a logical left shift of u bit-places on x_1 , giving x_2 . Put zero into the bits vacated by the shift. If u is greater than or equal to the number of bits in a word, x_2 is zero.

rshift

$$x_1 \ u \rightarrow x_2$$

Perform a logical right shift of u bit-places on x_1 , giving x_2 . Put zero into the bits vacated by the shift. If u is greater than or equal to the number of bits in a word, x_2 is zero.

arshift

$$x_1 \ u \rightarrow x_2$$

Perform an arithmetic right shift of u bit-places on x_1 , giving x_2 . Copy the original most-significant bit into the bits vacated by the shift. If u is greater than or equal to the number of bits in a word, all the bits of x_2 are the same as the original most-significant bit.

8.7 Comparison

These instructions compare two numbers on the stack, returning a flag; see section 8.5):

eq

$$s_1 \ s_2 \rightarrow flag$$

flag is 1 if and only if s_1 is equal to s_2 .

lt

$$s_1 \ s_2 \rightarrow flag$$

flag is 1 if and only if s_1 is less than s_2 .

ult

$$u_1 \ u_2 \rightarrow flag$$

flag is 1 if and only if u_1 is less than u_2 .

8.8 Control

These instructions implement unconditional and conditional branches, and subroutine call and return, with and without error handling.

jump

$$[a-addr] \rightarrow$$

If *ir* is 0, the target address is $a-addr$, otherwise it is $ir \times \text{word_bytes} + pc$. Set *pc* to the target address, and *ir* to 0.

jumpz

$$flag \ [a-addr] \rightarrow$$

If *flag* is false then perform the action of **jump**; otherwise, set *ir* to 0.

call

$$x_{u_1} \dots x_0 \ u_1 \ u_2 \ [a-addr_1] \rightarrow u_2 \ a-addr_2 ; x_{u_1} \dots x_0$$

Let the initial value of *pc* be $a-addr_2$. Perform the action of **jump**. Move $x_{u_1} \dots x_0$ from the caller's stack to the callee's.

ret

$$u \ a-addr ; x_u \dots x_0 \rightarrow x_u \dots x_0 \\ \text{or} \quad u \ a-addr \mid x_u \dots x_0 \rightarrow x_u \dots x_0 \ 0$$

Pop the current computation stack from the call stack; if the call stack is now empty, pop the current call stack. Set *pc* to $a-addr$, and move $x_u \dots x_0$ to the new stack. Set *ir* to 0. If a call stack was popped, push 0 on top of the stack.

throw $u \text{ a-addr} \mid n \rightarrow n$

Throw error n . Pop the current call stack from the catch stack. Pop the return address and number of return items, and push the error code.

catch $x_u \dots x_0 \ u_1 \ u_2 \ \text{a-addr}_1 \rightarrow u_2 \ \text{a-addr}_2 \mid x_u \dots x_0$

Perform the action of call as if ir were 0, then push a new call stack on to the catch stack and move the top-most computation stack from the previous call stack to the new one.

8.9 Extra instructions

Extra instructions, using the extra instruction, offer necessary functionality too rare or slow to deserve a core instruction.

extra

Perform extra instruction ir ; if ir is not the code of a valid extra instruction, throw error -1 . Extra instruction code 0 is used for instruction fetch (see section 7). The stack effect depends on the extra instruction.

8.10 Traps

Traps, using the trap instruction, are similar to extra instructions, but intended to be implementable as add-ons to an implementation, rather than as an integrated part of it. Traps may modify the memory and stack, but may not directly change the values of registers.

trap

Perform trap ir ; if ir is not the code of a valid trap, throw error -1 . The stack effect depends on the trap. Trap code -1 is used for instruction fetch (see section 7).

8.11 Instruction encoding

Instructions are encoded as bytes, packed into words, which are executed as described in section 7. The bytes have the following internal structure:

| 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
|---------------------|---|---|---|---|---|---|---|
| instruction | | | | | 0 | 0 | 0 |
| pushi $n < 0$ | | | | | 1 | 0 | 0 |
| pushreli $n < 0$ | | | | | 1 | 0 | |
| pushreli $n \geq 0$ | | | | | 0 | 1 | |
| pushi $n \geq 0$ | | | | | 0 | 1 | 1 |
| instruction | | | | | 1 | 1 | 1 |

Table 2 lists the instructions whose least-significant 3 bits are 000 or 111. Other instruction opcodes with those endings are invalid. Table 3 lists the extra instruction opcodes.

| Opcode | Instruction | Opcode | Instruction |
|--------|-------------|--------|-------------|
| 0x00 | extra | 0x80 | load |
| 0x08 | not | 0x88 | store |
| 0x10 | and | 0x90 | load1 |
| 0x18 | or | 0x98 | store1 |
| 0x20 | xor | 0xa0 | load2 |
| 0x28 | lshift | 0xa8 | store2 |
| 0x30 | rshift | 0xb0 | load4 |
| 0x38 | arshift | 0xb8 | store4 |
| 0x40 | pop | 0xc0 | push |
| 0x48 | dup | 0xc8 | pushrel |
| 0x50 | set | 0xd0 | negate |
| 0x58 | swap | 0xd8 | add |
| 0x60 | jump | 0xe0 | mul |
| 0x68 | jumpz | 0xe8 | eq |
| 0x70 | call | 0xf0 | lt |
| 0x78 | ret | 0xf8 | ult |

Table 2: Instruction opcodes

| Opcode | Instruction |
|--------|-------------|
| 0x1 | divmod |
| 0x2 | udivmod |
| 0x3 | catch |
| 0x4 | throw |

Table 3: Extra instruction opcodes

9 External interface

- Implementations should provide an API to create and run virtual machine code, and to add traps.
- Implementations can add extra instructions to provide extra computational primitives and other deeply-integrated facilities, and traps to offer access to system facilities, native libraries and so on; see section 8.9.

Acknowledgements

Martin Richards introduced me to Cintcode [2], which kindled my interest in virtual machines, and led to Beetle [3] and Mite [4], of which Mit is a sort of synthesis. GNU *lightning* [1] helped inspire me to greater simplicity, while still aiming for speed. Alistair Turnbull has for many years been a fount of ideas and criticism for all my work in computation, and lately a staunch collaborator on Mit.

References

- [1] Paulo Bonzini. Using and porting GNU *lightning*, 2000. <ftp://alpha.gnu.org/gnu/>.
- [2] Martin Richards. Cintcode distribution, 2000. <https://www.cl.cam.ac.uk/~mr/BCPL.html>.
- [3] Reuben Thomas. Beetle and pForth: a Forth virtual machine and compiler. BA dissertation, University of Cambridge, 1995. <https://rrt.sc3d.org/>.
- [4] Reuben Thomas. *Mite: a basis for ubiquitous virtual machines*. PhD thesis, University of Cambridge Computer Laboratory, November 2000. <https://rrt.sc3d.org/>.