The SMite Virtual Machine

Reuben Thomas

13th September 2018

Typographical conventions

Instructions and SMite's registers are shown in `Typewriter` font; interface calls are shown in **Bold** type, and followed by empty parentheses.

Addresses are given in bytes and refer to SMite's address space except where stated. Addresses are written in hexadecimal; hex numbers are prefixed with "0x".

1 Introduction

SMite is a simple virtual machine for study and experiment. It is a stack machine, based on the more complex register machine [3]. This paper gives a full description of SMite, but certain implementation-dependent features, such as the size of the stacks, are purposely left unspecified, and the exact method of implementation is left to the implementor in many particulars.

SMite is self-contained. Machine code routines on the host computer may be accessed using the `CALL_NATIVE` instruction, which can also be used to implement I/O (see 3.10). SMite supports a simple object module format.

SMite is conceptually (and usually in fact) a library, embedded in other programs. A small interface is provided for other programs to control SMite.

2 Architecture

SMite's address unit is the byte, which is eight bits. Words are four bytes. The word is the size of the numbers and addresses on which SMite operates, and of the items placed on the stacks. The word size is fixed to ensure compatibility of object code between implementations on different machines; the size of the byte and word have been chosen with a view to making efficient implementation of SMite possible on the vast majority of current machine architectures.

Words may have the bytes stored in big-endian or little-endian order. The address of a word is that of the byte in it with the lowest address.

2.1 Registers

The registers, each with its function, are set out in table 1.

All of the registers are word quantities except for `I` and `ENDISM`, which are one byte.

To ease efficient implementation, the registers may only be accessed by instructions (see section 3.6); not all registers are accessible, and only a few are writable.

Register	Function
PC	The Program Counter. Points to the next byte from which an instruction may be loaded.
I	The Instruction. Holds the opcode of an instruction to be executed.
MEMORY	The size in bytes of SMite's main memory, which must be a whole number of words.
SP	The data Stack Pointer.
RP	The Return stack Pointer.
S0	The data Stack base.
R0	The Return stack base.
SSIZE	The number of words allocated for the data stack.
RSIZE	The number of words allocated for the return stack.
HANDLER	The address PC is set to when an exception is raised.
ENDISM	The endianness of SMite: 0 = Little-endian, 1 = Big-endian.
BADPC	The address of the instruction that raised the most recent exception.
INVALID	The last address which caused an address exception.
PSIZE	The number of words in a host machine pointer (for CALL_NATIVE; see section 3.10).

Table 1: SMite's registers

2.2 Memory

SMite's memory is a discontiguous sequence of bytes with addresses in the range 0 to $2^{32} - 1$. Some locations may be read-only. The memory is contiguous in the range 0 to MEMORY $- 1$; this part is referred to as "main memory".

2.3 Stacks

The data and return stacks are word-aligned LIFO stacks of words. The stack pointers point to the top stack item on each stack. To **push** an item on to a stack means to store the item in the word beyond the stack pointer and then adjust the pointer to point to it; to **pop** an item means to make the pointer point to the second item on the stack. Instructions that change the number of items on a stack implicitly pop their arguments and push their results.

The data stack is used for passing values to instructions and routines and the return stack for holding subroutine return addresses. The return stack may be used for other operations subject to the restrictions placed on it by its normal usage: it must be returned before a RET instruction to the state it was in directly after the corresponding CALL.

In what follows, for "the stack" read "the data stack"; the return stack is always mentioned explicitly.

2.4 Operation

Before SMite is started, ENDISM should be set to 0 or 1 according to the implementation, and PSIZE to the appropriate value. The other registers should be initialised to 0.

MEMORY, ENDISM and PSIZE must not change while SMite is executing.

SMite is started by a call to the interface calls **run()** or **single_step()** (see section 4.2). In the former case, the execution cycle is entered:

> begin
>> load I from the byte pointed to by PC
>> add 1 to PC
>> execute the instruction in I
> repeat

In the latter case, the contents of the execution cycle is executed once, and control returns to the calling program.

Note that the calls **run()** and **single_step()** do not perform the initialisation specified above.

2.5 Termination

When SMite encounters a HALT instruction (see section 3.9), it returns the top data stack item as the reason code, unless SP does not point to a valid word, in which case reason code -257 is returned (see section 2.6).

Reason codes which are also valid exception codes (either reserved (see section 2.6) or user exception codes) should not normally be used. This allows exception codes to be passed back by an exception handler to the calling program, so that the calling program can handle certain exceptions without confusing exception codes and reason codes.

2.6 Exceptions

Exceptional and error conditions may be dealt with by **raising** an **exception**. The type of exception is signalled by its **exception code**. Exception codes are signed numbers. -1 to -511 are reserved for SMite's own exception codes; the meanings of those that may be raised by SMite are shown in table 2.

Raising an exception proceeds as follows:

> push the exception code on to the stack
> set BADPC to PC -1
> set PC to the value of HANDLER

The instruction that caused the exception to be raised is then considered to have finished executing.

An exception can be raised explicitly by THROW (see section 3.9); exceptions are also raised by various errors, such as an attempt to access an invalid address, or divide by zero.

Code	Meaning
-9	Invalid address (see below).
-10	Division by zero attempted (see section 3.3).
-20	Attempt to write to a read-only memory location.
-23	Address alignment exception (see below).
-256	Invalid opcode (see section 3.11).

Table 2: Exceptions raised by SMite

Exception -9 is raised whenever an attempt is made to access an invalid address, either by an instruction, or during an instruction fetch (because `PC` contains an invalid address). Exception -23 is raised when an instruction expecting an address of type `a-addr` (word-aligned) is given a non-aligned address. When SMite raises an address exception (-9 or -23), `INVALID` is first set to the offending address.

If `SP` is unaligned when an exception is raised, or putting the code on the stack would cause `SP` to be out of range, the effect of a `HALT` with code -257 is performed (although the actual mechanics are not, as that too would involve putting a number on the stack).

3   Instruction set

The instruction set is listed in sections 3.1 to 3.10, with the instructions grouped according to function. The instructions are given in the following format:

```
NAME                                    ( before -- after )
                                   R: ( before -- after )
```
 Description.

The first line consists of the name of the instruction. On the right are the stack effect or effects, which show the effect of the instruction on the data and return (R) stacks. Underneath is the description.

Stack effects are written
$$(\textit{before} -- \textit{after})$$
where *before* and *after* are stack pictures showing the items on top of a stack before and after the instruction is executed. An instruction only affects the items shown in its stack effects. The brackets and dashes serve merely to delimit the stack effect and to separate *before* from *after*. **Stack pictures** are a representation of the top-most items on the stack, and are written
$$i_1 \ i_2 \ldots i_{n-1} \ i_n$$
where the i_k are stack items, each of which occupies a whole number of words, with i_n being on top of the stack. The symbols denoting different types of stack item are shown in table 3.

Symbol	Data type
flag	flag
true	true flag
false	false flag
byte	byte
n	signed number
u	unsigned number
n|u	number (signed or unsigned)
x	unspecified word
addr	address
a-addr	word-aligned address

Table 3: Types used in stack effects

Types are only used to indicate how instructions treat their arguments and results; SMite does not distinguish between stack items of different types. In stack pictures the most general argument types with which each instruction can be supplied are given; subtypes may be substituted. Using the phrase "$i \Rightarrow j$" to

4

denote "i is a subtype of j", table 4 shows the subtype relationships. The subtype relation is transitive.

$u \Rightarrow x$
$n \Rightarrow x$
$byte \Rightarrow u$
$a\text{-}addr \Rightarrow addr \Rightarrow u$
$flag \Rightarrow x$

Table 4: The subtype relation

Numbers are represented in twos complement form. `addr` consists of all valid addresses. Numeric constants can be included in stack pictures, and are of type `n|u`.

Each type may be suffixed by a number in stack pictures; if the same combination of type and suffix appears more than once in a stack effect, it refers to identical stack items. Alternative `after` pictures are separated by "`|`", and the circumstances under which each occurs are detailed in the instruction description.

The symbols `i*x`, `j*x` and `k*x` are used to denote different collections of zero or more words of any data type. Ellipsis is used for indeterminate numbers of specified types of word.

If an instruction does not modify the return stack, the corresponding stack picture is omitted.

3.1 Stack manipulation

These instructions manage the data stack and move values between stacks.

POP $\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ ($x_u \ldots x_1$ u `--`)

Remove u items from the stack.

PUSH $\qquad\qquad\qquad\qquad\qquad$ ($x_u \ldots x_0$ u `--` $x_u \ldots x_0$ x_u)

Remove u. Copy x_u to the top of the stack.

SWAP $\qquad\qquad\qquad\qquad$ ($x_u \ldots x_0$ u `--` $x_0 x_{u-1} \ldots x_1$ x_u)

Exchange the top stack item with the uth. If u is zero, do nothing.

POP2R $\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ (x `--`)
$\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ R: (`--` x)

Move x to the return stack.

RPOP $\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ (`--` x)
$\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ R: (x `--`)

Move x from the return stack to the data stack.

RPUSH $\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ (u `--` x_u)
$\qquad\qquad\qquad$ R: (x_u $x_{u-1} \ldots x_0$ `--` x_u $x_{u-1} \ldots x_0$)

Remove u. Copy x_u to the top of the data stack.

3.2 Comparison

These words compare two numbers (or, for equality tests, any two words) on the stack, returning a flag, true with all bits set if the test succeeds and false otherwise.

LT $\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ (n_1 n_2 `--` `flag`)

`flag` is true if and only if n_1 is less than n_2.

EQ   (x_1 x_2 -- *flag*)

flag is true if and only if x_1 is bit-for-bit the same as x_2.

ULT   (u_1 u_2 -- *flag*)

flag is true if and only if u_1 is less than u_2.

3.3 Arithmetic

These instructions consist of monadic and dyadic operators. All calculations are made without bounds or overflow checking, except as detailed for certain instructions.

Addition:

ADD   ($n_1|u_1$ $n_2|u_2$ -- $n_3|u_3$)

Add $n_2|u_2$ to $n_1|u_1$, giving the sum $n_3|u_3$.

Multiplication and division (note that all division instructions raise exception -10 if division by zero is attempted):

MUL   ($n_1|u_1$ $n_2|u_2$ -- $n_3|u_3$)

Multiply $n_1|u_1$ by $n_2|u_2$ giving the product $n_3|u_3$.

UDIVMOD   (u_1 u_2 -- u_3 u_4)

Divide u_1 by u_2, giving the single-word quotient u_3 and the single-word remainder u_4.

DIVMOD   (n_1 n_2 -- n_3 n_4)

Divide n_1 by n_2 using symmetric division, giving the single-word quotient n_3 and the single-word remainder n_4. The quotient is rounded towards zero.

Sign function:

NEGATE   (n_1 -- n_2)

Negate n_1, giving its arithmetic inverse n_2.

3.4 Logic and shifts

These instructions consist of bitwise logical operators and bitwise shifts.

Logic functions:

INVERT   (x_1 -- x_2)

Invert all bits of x_1, giving its logical inverse x_2.

AND   (x_1 x_2 -- x_3)

x_3 is the bit-by-bit logical "and" of x_1 with x_2.

OR   (x_1 x_2 -- x_3)

x_3 is the bit-by-bit inclusive-or of x_1 with x_2.

XOR   (x_1 x_2 -- x_3)

x_3 is the bit-by-bit exclusive-or of x_1 with x_2.

Shifts:

LSHIFT   (x_1 u -- x_2)

Perform a logical left shift of u bit-places on x_1, giving x_2. Put zero into the least significant bits vacated by the shift. If u is greater than or equal to 32, x_2 is zero.

RSHIFT   (x_1 u -- x_2)

Perform a logical right shift of u bit-places on x_1, giving x_2. Put zero into the most significant bits vacated by the shift. If u is greater than or equal to 32, x_2 is zero.

3.5 Memory

These instructions fetch and store words and bytes to and from memory; there is also an instruction to add a number to another stored in memory.

`LOAD`	(*a-addr -- x*)

x is the value stored at *a-addr*.

`STORE`	(*x a-addr --*)

Store *x* at *a-addr*.

`LOADB`	(*addr -- byte*)

If `ENDISM` is 1, exclusive-or *addr* with 3. Fetch the byte stored at *addr*. The unused high-order bits are all zeroes.

`STOREB`	(*byte addr --*)

If `ENDISM` is 1, exclusive-or *addr* with 3. Store *byte* at *addr*. Only one byte is transferred.

3.6 Registers

As mentioned in section 2.1, the stack pointers SP and RP may only be accessed through special instructions:

`PUSH_SP`   (*-- a-addr*)

a-addr is the value of SP.

`STORE_SP`   (*a-addr --*)

Set SP to *a-addr*.

`PUSH_RP`   (*-- a-addr*)

a-addr is the value of RP.

`STORE_RP`   (*a-addr --*)

Set RP to *a-addr*.

`PUSH_PC`   (*-- addr*)

Push PC on to the stack.

`PUSH_S0`   (*-- a-addr*)

Push S0 on to the stack.

`PUSH_SSIZE`   (*-- u*)

Push SSIZE on to the stack.

`PUSH_R0`   (*-- a-addr*)

Push R0 on to the stack.

`PUSH_RSIZE`   (*-- u*)

Push RSIZE on to the stack.

`PUSH_HANDLER`   (*-- addr*)

Push HANDLER on to the stack.

`STORE_HANDLER`   (*addr --*)

Set HANDLER to *addr*.

`PUSH_MEMORY`   (*-- a-addr*)

Push MEMORY on to the stack.

`PUSH_BADPC`   (*-- addr*)

Push BADPC on to the stack.

```
PUSH_INVALID                                              ( -- addr )
```
Push `INVALID` on to the stack.

```
PUSH_PSIZE                                                ( -- u )
```
u is the value of `PSIZE`.

3.7 Control structures

These instructions implement unconditional and conditional branches, and sub-routine call and return; there is also a no-op.

No-op:

```
NOP                                                       ( -- )
```
Do nothing.

Branches:

```
BRANCH                                                    ( addr -- )
```
Set PC to *addr*.

```
BRANCHZ                                                   ( flag addr -- )
```
If *flag* is false then set PC to *addr*.

Subroutine call and return:

```
CALL                                                      ( a-addr₁ -- )
                                                     R: (  -- a-addr₂ )
```
Push PC on to the return stack, put $a\text{-}addr_1$ into PC.

```
RET                                                       ( -- )
                                                     R: ( addr -- )
```
Put *addr* into PC.

3.8 Literal numbers

Literal numbers are word quantities encoded by one or more bytes, as follows: the significant bits of the number are split into groups of six bits, starting at the least significant end. The chunks are stored in consecutive bytes. All but the last byte have the seventh bit set and eighth bit clear; the final byte has the top two bits either both set or both clear, to match the number's most significant bit.

When a literal is executed, it is pushed on to the stack, and execution resumes at the first byte that did not form part of the number.

3.9 Exceptions

These instructions give access to SMite's exception mechanisms (see section 2.6).

```
THROW                                                     ( n -- )
```
Raise exception *n*.

```
HALT                                                      ( x -- )
```
Stop SMite, returning reason code *x* to the calling program (see section 4.2). If SP is out of range or unaligned, -257 is returned as the reason code.

3.10 External access

These instructions allow access to SMite's libraries, the operating system and native machine code.

`CALL_NATIVE`   (*i*x -- *)

Make a subroutine call to the routine at the address given (in the host machine's format, padded out to a number of words) on the data stack. The size and format of this address are machine-dependent.

`EXTRA`   (*i*x -- j*x *)

Perform implementation-dependent actions; for example, this can be used to implement system-dependent functionality such as I/O.

3.11 Opcodes

Table 5 lists the opcodes in numerical order. Undefined opcodes raise exception −256. Opcodes 0x00–0x7f and 0xc0–0xff are used for literal numbers (see section 3.8). Other opcodes are undefined.

Opcode	Instruction	Opcode	Instruction	Opcode	Instruction
0x80	NOP	0x90	AND	0xa0	EXTRA
0x81	POP	0x91	OR	0xa1	PUSH_PSIZE
0x82	PUSH	0x92	XOR	0xa2	PUSH_SP
0x83	SWAP	0x93	LSHIFT	0xa3	STORE_SP
0x84	RPUSH	0x94	RSHIFT	0xa4	PUSH_RP
0x85	POP2R	0x95	LOAD	0xa5	STORE_RP
0x86	RPOP	0x96	STORE	0xa6	PUSH_PC
0x87	LT	0x97	LOADB	0xa7	PUSH_S0
0x88	EQ	0x98	STOREB	0xa8	PUSH_SSIZE
0x89	ULT	0x99	BRANCH	0xa9	PUSH_R0
0x8a	ADD	0x9a	BRANCHZ	0xaa	PUSH_RSIZE
0x8b	MUL	0x9b	CALL	0xab	PUSH_HANDLER
0x8c	UDIVMOD	0x9c	RET	0xac	STORE_HANDLER
0x8d	DIVMOD	0x9d	THROW	0xad	PUSH_MEMORY
0x8e	NEGATE	0x9e	HALT	0xae	PUSH_BADPC
0x8f	INVERT	0x9f	CALL_NATIVE	0xaf	PUSH_INVALID

Table 5: SMite's opcodes

4 External interface

SMite's external interface comes in three parts. The calling interface allows SMite to be controlled by other programs. The `CALL_NATIVE` instruction (see section 3.10) allows implementations to provide access to system facilities, previously written code, code written in other languages, and the speed of machine code in time-critical situations. The object module format allows compiled code to be saved, reloaded and shared between systems.

4.1 Object module format

The first seven bytes of an object module should be the ASCII codes of the letters "smite" padded with ASCII NULs (0x00), then the one-byte contents of the `ENDISM`

register of the system which saved the module. The next word should contain the number of bytes the code occupies. The number must have the same endianness as that indicated in the previous byte. Then follows the code.

Object modules have a simple structure, as they are only intended for loading an initial memory image into SMite.

4.2   Calling interface

The calling interface is difficult to specify with the same precision as the rest of SMite, as it may be implemented in any language. However, since only basic types are used, and the semantics are simple, it is expected that implementations in different language producing the same result will be easy to program. A Modula-like syntax is used to give the definitions here. Implementation-defined error codes must be documented, but are optional. All addresses passed as parameters must be word-aligned. A SMite must provide the following calls:

native_address (*integer, boolean*) : pointer
> Return a native pointer corresponding to the given SMite address. If the SMite address is invalid, or the Boolean flag is true and the address is read-only, then a distinguished invalid pointer is returned.

run () : integer
> Start SMite by entering the execution cycle as described in section 2.4. If SMite ever executes a `HALT` instruction (see section 3.9), the reason code is returned as the result.

single_step () : integer
> Execute a single pass of the execution cycle, and return reason code -259, unless a `HALT` instruction was obeyed (see section 3.9), in which case the reason code passed to it is returned.

load_object (*file, address*) : integer
> Load the object module specified by *file*, which may be a filename or some other specifier, to the SMite address *address*. First the module's header is checked; if the first seven bytes are not as specified above in section 4.1, or the endianness value is not 0 or 1, then return -2. If the code will not fit into memory at the address given, or the address is out of range or unaligned, return -1. Otherwise load the code into memory, converting it if the endianness value is different from the current value of `ENDISM`. The result is 0 if successful, and some other implementation-defined value if there is a filing system or other error.

SMite must also provide access to its registers and address space through appropriate data objects.

Acknowledgements

Martin Richards's demonstration of his BCPL-oriented Cintcode virtual machine [1] convinced me it was going to be fun working on virtual machines. He also supervised my BA dissertation project, Beetle [2], on which SMite is based.

References

[1] Martin Richards. Cintcode distribution, 2000. `https://www.cl.cam.ac.uk/~mr/BCPL.html`.

[2] Reuben Thomas. Beetle and pForth: a Forth virtual machine and compiler. BA dissertation, University of Cambridge, 1995. `https://rrt.sc3d.org/`.

[3] Reuben Thomas. *Mite: a basis for ubiquitous virtual machines*. PhD thesis, University of Cambridge Computer Laboratory, November 2000. `https://rrt.sc3d.org/`.